

BeSafe Hackathon

January 11, 2026

```
[1]: import pandas as pd
import numpy as np
import xgboost as xgb
import matplotlib.pyplot as plt
import kagglehub
import os
from sklearn.model_selection import train_test_split, KFold, cross_val_score
from sklearn.metrics import accuracy_score

# 1. Load Data
# -----
if 'df' not in locals():
    print("Downloading dataset...")
    path = kagglehub.dataset_download("sid321axn/malicious-urls-dataset")
    csv_file = [f for f in os.listdir(path) if f.endswith('.csv')][0]
    full_path = os.path.join(path, csv_file)
    print(f"Loading CSV from: {full_path}")
    df = pd.read_csv(full_path)

# 2. Encode Labels
# -----
# 0 = Safe (benign), 1 = Malicious (defacement, phishing, malware)
df['label'] = df['type'].apply(lambda x: 0 if x == 'benign' else 1)

# 3. Visualize Data Balance (Critical Step)
# -----
label_counts = df['label'].value_counts()
print("Data Distribution:\n", label_counts)

plt.figure(figsize=(8, 5))
# Plotting bar chart of Benign vs Malicious
colors = ['#4CAF50', '#F44336'] # Green for Safe, Red for Bad
plt.bar(['Safe (0)', 'Malicious (1)'], label_counts.values, color=colors)
plt.title('Dataset Balance: Safe vs Malicious URLs', fontsize=14)
plt.ylabel('Number of Samples')
plt.grid(axis='y', linestyle='--', alpha=0.7)
```

```
# Add text on top of bars
for i, v in enumerate(label_counts.values):
    plt.text(i, v + 500, str(v), ha='center', fontweight='bold')

plt.show()

# Planner Note: If the red bar is much lower than the green bar,
# we have an "Imbalanced Dataset". XGBoost has a fix for this
# called 'scale_pos_weight', which we will use later.
```

C:\Users\lina3\AppData\Local\Programs\Python\Python313\Lib\site-packages\tqdm\auto.py:21: TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html

```
from .autonotebook import tqdm as notebook_tqdm
```


Downloading dataset...

Warning: Looks like you're using an outdated `kagglehub` version (installed: 0.3.13), please consider upgrading to the latest version (0.4.0).

Loading CSV from: C:\Users\lina3\.cache\kagglehub\datasets\sid321axn\malicious-urls-dataset\versions\1\malicious_phish.csv

Data Distribution:

```
label
0    428103
1    223088
Name: count, dtype: int64
```



```
[2]: print("--- Data Overview ---")
print(f"Total Rows: {len(df)}")
print("\n--- Labels Distribution (Original) ---")
print(df['type'].value_counts())

print("\n--- Labels Distribution (Percentage) ---")
print(df['type'].value_counts(normalize=True) * 100)

print("\n--- Sample Data (First 5 rows) ---")
display(df.head()) # display Jupyter-
```

--- Data Overview ---

Total Rows: 651191

--- Labels Distribution (Original) ---

type

benign 428103

defacement 96457

phishing 94111

malware 32520

Name: count, dtype: int64

--- Labels Distribution (Percentage) ---

type

benign 65.741541

defacement 14.812398

phishing 14.452135

malware 4.993927

Name: proportion, dtype: float64

--- Sample Data (First 5 rows) ---

	url	type	label
0	br-icloud.com.br	phishing	1
1	mp3raid.com/music/krizz_kaliko.html	benign	0
2	bopsecrets.org/rexroth/cr/1.htm	benign	0
3	http://www.garage-pirenne.be/index.php?option=...	defacement	1
4	http://adventure-nicaragua.net/index.php?optio...	defacement	1

```
[3]: import re
from urllib.parse import urlparse

def advanced_feature_extraction(url):
    url = str(url)
    features = []
```

```

# --- Structural Features (The Shape of the URL) ---

# 1. Total Length
features.append(len(url))

# 2. Count Dots (Phishing often uses many subdomains: a.b.c.bank.com)
features.append(url.count('.'))

# 3. Count Hyphens (Legit sites rarely use: secure-login-account-update)
features.append(url.count('-'))

# 4. Count Special Obfuscation Chars (@, %, // in middle)
features.append(url.count('@'))
features.append(url.count('%'))

# 5. Directory Depth (How many slashes after the domain?)
# http://com/a/b/c/d/e -> depth 5
features.append(url.count('/'))

# --- Statistical Features (Math on the string) ---

# 6. Digit to Letter Ratio (Malware links often have random numbers)
digits = sum(c.isdigit() for c in url)
letters = sum(c.isalpha() for c in url)
features.append(digits / (letters + 1)) # +1 to avoid division by zero

# 7. Has IP Address? (e.g. http://192.168.1.1/login)
# Using a simple Regex to find IP patterns
has_ip = 1 if re.search(r'\d{1,3}\.\d{1,3}\.\d{1,3}\.\d{1,3}', url) else 0
features.append(has_ip)

# --- Keyword Features (The "Intent") ---

# 8. Suspicious Keywords (Financial/Urgency)
suspicious_words = ['login', 'signin', 'bank', 'secure', 'account',
↪ 'update',
                    'verify', 'confirm', 'wallet', 'crypto', 'bonus',
↪ 'free']
features.append(sum(1 for word in suspicious_words if word in url.lower()))

# 9. Brand Spoofing (Commonly targeted brands)
brands = ['paypal', 'apple', 'google', 'microsoft', 'facebook', 'netflix',
↪ 'amazon']
features.append(sum(1 for brand in brands if brand in url.lower()))

# 10. File Extensions (Is it trying to download something?)
dangerous_ext = ['.exe', '.zip', '.tar', '.js', '.vbs', '.apk']

```

```

        features.append(1 if any(ext in url.lower() for ext in dangerous_ext) else 0)

    # 11. Is HTTPS? (0 if http, 1 if https - though many phishers use https now)
    features.append(1 if 'https' in url.lower() else 0)

    # 12. "www" presence (Random generated domains often skip 'www')
    features.append(1 if 'www.' in url.lower() else 0)

    return features

# Apply to DataFrame
print("Extracting features... (This might take a minute)")
X_temp = df['url'].apply(lambda x: advanced_feature_extraction(x))

# Convert list of lists to DataFrame
feature_names = [
    'len', 'dots', 'hyphens', 'at_symbol', 'percent', 'slashes',
    'digit_ratio', 'has_ip', 'sus_keywords', 'brands', 'is_file', 'is_https',
    'has_www'
]
X = pd.DataFrame(X_temp.tolist(), columns=feature_names)
y = df['label']

print("Feature Matrix Shape:", X.shape)
print("First 5 rows of features:")
display(X.head()) # Or print(X.head()) if not in Jupyter

```

Extracting features... (This might take a minute)

Feature Matrix Shape: (651191, 13)

First 5 rows of features:

	len	dots	hyphens	at_symbol	percent	slashes	digit_ratio	has_ip	\
0	16	2	1	0	0	0	0.000000	0	
1	35	2	0	0	0	2	0.033333	0	
2	31	2	0	0	0	3	0.038462	0	
3	88	3	1	0	0	3	0.109375	0	
4	235	2	1	0	0	3	0.110000	0	

	sus_keywords	brands	is_file	is_https	has_www
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	0	0	0	1
4	0	0	0	0	0

```

[14]: import seaborn as sns
import pandas as pd

```

```

import matplotlib.pyplot as plt
import xgboost as xgb
from sklearn.model_selection import cross_validate # <--- The new tool

# 1. Search Ranges
param_grid = {
    'max_depth': [10, 15, 20, 25, 30, 35],
    'learning_rate': [0.01, 0.05, 0.1, 0.15, 0.2]
}

results = []

print("Starting Deep Dive Grid Search (Train vs Test)...")
print("-----")
print(f"{'Depth':<6} | {'Rate':<6} | {'Train Acc':<10} | {'Test Acc':<10} | ␣
    ↳{'Gap (Overfit)':<12}")
print("-----")

# 2. Double Loop
for depth in param_grid['max_depth']:
    for rate in param_grid['learning_rate']:

        model = xgb.XGBClassifier(
            n_estimators=100,
            max_depth=depth,
            learning_rate=rate,
            eval_metric='logloss'
        )

        # Use cross_validate to get BOTH scores
        cv_results = cross_validate(
            model,
            X_train,
            y_train,
            cv=3,
            scoring='accuracy',
            return_train_score=True # <--- This is the magic key
        )

        # Calculate averages
        mean_train = cv_results['train_score'].mean()
        mean_test = cv_results['test_score'].mean()
        gap = mean_train - mean_test

        results.append([depth, rate, mean_train, mean_test, gap])

# Print row-by-row

```

```

        print(f"{depth:<6} | {rate:<6} | {mean_train*100:.2f}% | \n
↳ {mean_test*100:.2f}% | {gap*100:.2f}%")

print("-----")

# 3. Visualize Results
df_results = pd.DataFrame(results, columns=['Max Depth', 'Learning Rate', \n
↳ 'Train Accuracy', 'Test Accuracy', 'Gap'])

# Create two heatmaps side-by-side
fig, ax = plt.subplots(1, 2, figsize=(16, 6))

# Heatmap 1: Test Accuracy (What we want to maximize)
pivot_test = df_results.pivot(index='Max Depth', columns='Learning Rate', \n
↳ values='Test Accuracy')
sns.heatmap(pivot_test, annot=True, fmt=".4f", cmap='viridis', ax=ax[0])
ax[0].set_title('TEST Accuracy (Higher is Better)')

# Heatmap 2: Overfitting Gap (What we want to minimize)
pivot_gap = df_results.pivot(index='Max Depth', columns='Learning Rate', \n
↳ values='Gap')
sns.heatmap(pivot_gap, annot=True, fmt=".4f", cmap='coolwarm', ax=ax[1])
ax[1].set_title('Overfitting GAP (Lower is Better)')

plt.show()



# 4. Auto-select the winner (Based on best TEST accuracy)
best_row = df_results.loc[df_results['Test Accuracy'].idxmax()]
print("\n WINNER PARAMETERS:")
print(f"Best Depth: {int(best_row['Max Depth'])}")
print(f"Best Learning Rate: {best_row['Learning Rate']}")
print(f"Test Accuracy: {best_row['Test Accuracy']*100:.2f}%")
print(f"Train Accuracy: {best_row['Train Accuracy']*100:.2f}% (Gap: \n
↳ {best_row['Gap']*100:.2f}%)")

```

Starting Deep Dive Grid Search (Train vs Test)...

Depth	Rate	Train Acc	Test Acc	Gap (Overfit)
10	0.01	90.50%	90.36%	0.14%
10	0.05	92.38%	92.08%	0.29%
10	0.1	93.23%	92.76%	0.47%
10	0.15	93.63%	93.04%	0.59%
10	0.2	93.98%	93.23%	0.75%
15	0.01	92.43%	91.99%	0.44%
15	0.05	93.76%	92.96%	0.80%
15	0.1	94.37%	93.32%	1.05%

15	0.15	94.86%	93.54%	1.31%
15	0.2	95.13%	93.63%	1.50%
20	0.01	93.73%	92.83%	0.89%
20	0.05	95.00%	93.52%	1.48%
20	0.1	95.35%	93.63%	1.72%
20	0.15	95.58%	93.70%	1.88%
20	0.2	95.76%	93.72%	2.04%
25	0.01	94.36%	93.21%	1.16%
25	0.05	95.57%	93.62%	1.94%
25	0.1	95.79%	93.68%	2.11%
25	0.15	95.95%	93.71%	2.24%
25	0.2	96.08%	93.71%	2.38%
30	0.01	94.49%	93.27%	1.22%
30	0.05	95.76%	93.64%	2.13%
30	0.1	96.01%	93.68%	2.32%
30	0.15	96.12%	93.69%	2.43%
30	0.2	96.24%	93.69%	2.55%
35	0.01	94.50%	93.27%	1.23%
35	0.05	95.80%	93.63%	2.17%
35	0.1	96.09%	93.67%	2.42%
35	0.15	96.23%	93.69%	2.54%
35	0.2	96.33%	93.68%	2.65%



WINNER PARAMETERS:

Best Depth: 20

Best Learning Rate: 0.2

Test Accuracy: 93.72%

Train Accuracy: 95.76% (Gap: 2.04%)


```

[15]: import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
import xgboost as xgb
from sklearn.model_selection import cross_validate # <--- The new tool

# 1. Search Ranges
param_grid = {
    'max_depth': [20],
    'learning_rate': [0.01, 0.05, 0.1, 0.15, 0.2, 0.25, 0.3, 0.4]
}

results = []

print("Starting Deep Dive Grid Search (Train vs Test)...")
print("-----")
print(f"{'Depth':<6} | {'Rate':<6} | {'Train Acc':<10} | {'Test Acc':<10} | ␣
    ↳{'Gap (Overfit)':<12}")
print("-----")

# 2. Double Loop
for depth in param_grid['max_depth']:
    for rate in param_grid['learning_rate']:

        model = xgb.XGBClassifier(
            n_estimators=100,
            max_depth=depth,
            learning_rate=rate,
            eval_metric='logloss'
        )

        # Use cross_validate to get BOTH scores
        cv_results = cross_validate(
            model,
            X_train,
            y_train,
            cv=3,
            scoring='accuracy',
            return_train_score=True # <--- This is the magic key
        )

        # Calculate averages
        mean_train = cv_results['train_score'].mean()
        mean_test = cv_results['test_score'].mean()
        gap = mean_train - mean_test

        results.append([depth, rate, mean_train, mean_test, gap])

```

```

        # Print row-by-row
        print(f"{depth:<6} | {rate:<6} | {mean_train*100:.2f}%      |  

↳ {mean_test*100:.2f}%      | {gap*100:.2f}%")

print("-----")

# 3. Visualize Results
df_results = pd.DataFrame(results, columns=['Max Depth', 'Learning Rate',  

↳ 'Train Accuracy', 'Test Accuracy', 'Gap'])

# Create two heatmaps side-by-side
fig, ax = plt.subplots(1, 2, figsize=(16, 6))

# Heatmap 1: Test Accuracy (What we want to maximize)
pivot_test = df_results.pivot(index='Max Depth', columns='Learning Rate',  

↳ values='Test Accuracy')
sns.heatmap(pivot_test, annot=True, fmt=".4f", cmap='viridis', ax=ax[0])
ax[0].set_title('TEST Accuracy (Higher is Better)')

# Heatmap 2: Overfitting Gap (What we want to minimize)
pivot_gap = df_results.pivot(index='Max Depth', columns='Learning Rate',  


↳ values='Gap')
sns.heatmap(pivot_gap, annot=True, fmt=".4f", cmap='coolwarm', ax=ax[1])
ax[1].set_title('Overfitting GAP (Lower is Better)')

plt.show()

```

Starting Deep Dive Grid Search (Train vs Test)...

Depth	Rate	Train Acc	Test Acc	Gap (Overfit)
20	0.01	93.73%	92.83%	0.89%
20	0.05	95.00%	93.52%	1.48%
20	0.1	95.35%	93.63%	1.72%
20	0.15	95.58%	93.70%	1.88%
20	0.2	95.76%	93.72%	2.04%
20	0.25	95.94%	93.72%	2.22%
20	0.3	96.03%	93.72%	2.31%
20	0.4	96.21%	93.69%	2.52%



WINNER PARAMETERS:

Best Depth: 20

Best Learning Rate: 0.3

Test Accuracy: 93.72%

Train Accuracy: 96.03% (Gap: 2.31%)

```
[16]: # 1. Retrieve the Winning Parameters from the previous step
# We access the 'best_row' variable we created in the Grid Search cell
final_depth = int(best_row['Max Depth'])
final_rate = best_row['Learning Rate']

print(f" Initializing Final Model with Winner Params: Depth={final_depth},
      ↪Rate={final_rate}")

# 2. Train the Final Model on ALL training data
final_model = xgb.XGBClassifier(
    n_estimators=200,          # We assume 200 is good for final polish
    max_depth=final_depth,    # The Winner
    learning_rate=final_rate, # The Winner
    eval_metric='logloss'
)

final_model.fit(X_train, y_train)

# 3. Final Evaluation on Test Set
# This is the "Moment of Truth" - checking on data the model has NEVER seen
y_pred_final = final_model.predict(X_test)
final_acc = accuracy_score(y_test, y_pred_final)

print(f"\n=====")
```

```

print(f" FINAL MODEL ACCURACY: {final_acc * 100:.2f}%")
print(f"=====")

# 4. Feature Importance
# See which new features (like 'dots' or 'keywords') actually mattered
plt.figure(figsize=(10, 8))
xgb.plot_importance(final_model, max_num_features=15, height=0.5,
    ↪importance_type='weight', title='Feature Importance (What signals phishing?
    ↪)')
plt.show()

# 5. Save the Brain
final_model.save_model("enhanced_url_classifier.json")
print("Model saved successfully.")

```


Initializing Final Model with Winner Params: Depth=20, Rate=0.3

```

=====
FINAL MODEL ACCURACY: 93.91%
=====

```

<Figure size 1000x800 with 0 Axes>



Model saved successfully.

Take 2

```
[ ]: import math
from collections import Counter
from urllib.parse import urlparse

# " "
SUSPICIOUS_TLDS = ['.xyz', '.top', '.club', '.win', '.gq', '.cn', '.tk', '.ml',
                  '.ga', '.cf', '.buzz']

# ) (
SHORTENERS = ['bit.ly', 'goo.gl', 'tinyurl.com', 'is.gd', 'cli.gs', 't.co']

def calculate_entropy(text):
    """
    Calculate the entropy of a string.
    """
    if not text:
        return 0
    entropy = 0
    total_len = len(text)
    #
    counts = Counter(text)

    for count in counts.values():
        p = count / total_len
        entropy -= p * math.log2(p)
    return entropy

def advanced_feature_extraction(url):
    url = str(url).lower() #
    features = []

    # --- ' ) ( ---
    features.append(len(url))
    features.append(url.count('.'))
    features.append(url.count('-'))
    features.append(url.count('@'))
    features.append(url.count('%'))
    features.append(url.count('/'))

    # --- ' ---

    # 1. ) , (
    digits = sum(c.isdigit() for c in url)
    letters = sum(c.isalpha() for c in url)
```

```

features.append(digits / (letters + 1))

# 2. NEW:          )      (
#                  ,
try:
    domain = urlparse(url).netloc
    if not domain: domain = url
except:
    domain = url
features.append(calculate_entropy(domain))

# ---      '          ---

# 3. NEW:          ? (TLD)
is_sus_tld = 1 if any(url.endswith(tld) for tld in SUSPICIOUS_TLDS) else 0
features.append(is_sus_tld)

# 4. NEW:          ?
is_shortener = 1 if any(s in url for s in SHORTENERS) else 0
features.append(is_shortener)

#          (      )
suspicious_words = ['login', 'signin', 'bank', 'secure', 'account', '
↪'update',
                    'verify', 'confirm', 'wallet', 'crypto', 'bonus', '
↪'free', 'ebay']
features.append(sum(1 for word in suspicious_words if word in url))

#          (      )
brands = ['paypal', 'apple', 'google', 'microsoft', 'facebook', 'netflix', '
↪'amazon']
features.append(sum(1 for brand in brands if brand in url))

#          (      )
dangerous_ext = ['.exe', '.zip', '.tar', '.js', '.vbs', '.apk']
features.append(1 if any(ext in url for ext in dangerous_ext) else 0)

# HTTPS WWW-      ),          ,      (
features.append(1 if 'https' in url else 0)
features.append(1 if 'www.' in url else 0)

return features

print("New Feature Extraction Logic Loaded! ")

```

```
[19]: print(" Applying NEW feature extraction logic to the dataset...")
```

```

print("    (Calculating Entropy & checking TLDs for all URLs - this takes a
↳moment)")

# 1.                                URL-
X_new_features = df['url'].apply(lambda x: pd.
↳Series(advanced_feature_extraction(x)))

# 2.                                )          (
feature_names = [
    'len', 'dots', 'hyphens', 'at_symbol', 'percent', 'slashes',
    'digit_ratio',          #
    'entropy',              # <---      !
    'is_sus_tld',           # <---      !
    'is_shortener',         # <---      !
    'sus_keywords', 'brands', 'dangerous_ext', 'is_https', 'has_www' #
]

X = pd.DataFrame(X_new_features.values, columns=feature_names)
y = df['label'] # Y-

# 3.      Train- Test- )      '      (
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

print(" Data updated! Now the model can see the Entropy and TLDs.")
print(f"    New Feature Matrix Shape: {X.shape}")

```

Applying NEW feature extraction logic to the dataset...

(Calculating Entropy & checking TLDs for all URLs - this takes a moment)

Data updated! Now the model can see the Entropy and TLDs.

New Feature Matrix Shape: (651191, 15)

```

[20]: import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
import xgboost as xgb
from sklearn.model_selection import cross_validate # <--- The new tool

# 1. Search Ranges
param_grid = {
    'max_depth': [10, 15, 20, 25, 30, 35],
    'learning_rate': [0.01, 0.05, 0.1, 0.15, 0.2]
}

results = []

print("Starting Deep Dive Grid Search (Train vs Test)...")

```

```

print("-----")
print(f"{'Depth':<6} | {'Rate':<6} | {'Train Acc':<10} | {'Test Acc':<10} | ␣
↪{'Gap (Overfit)':<12}")
print("-----")

# 2. Double Loop
for depth in param_grid['max_depth']:
    for rate in param_grid['learning_rate']:

        model = xgb.XGBClassifier(
            n_estimators=100,
            max_depth=depth,
            learning_rate=rate,
            eval_metric='logloss'
        )

        # Use cross_validate to get BOTH scores
        cv_results = cross_validate(
            model,
            X_train,
            y_train,
            cv=3,
            scoring='accuracy',
            return_train_score=True # <--- This is the magic key
        )

        # Calculate averages
        mean_train = cv_results['train_score'].mean()
        mean_test = cv_results['test_score'].mean()
        gap = mean_train - mean_test

        results.append([depth, rate, mean_train, mean_test, gap])

        # Print row-by-row
        print(f"{'depth':<6} | {'rate':<6} | {'mean_train*100:.2f}%      | ␣
↪{'mean_test*100:.2f}%      | {'gap*100:.2f}%"")

print("-----")

# 3. Visualize Results
df_results = pd.DataFrame(results, columns=['Max Depth', 'Learning Rate', ␣
↪'Train Accuracy', 'Test Accuracy', 'Gap'])

# Create two heatmaps side-by-side
fig, ax = plt.subplots(1, 2, figsize=(16, 6))

# Heatmap 1: Test Accuracy (What we want to maximize)

```



```

pivot_test = df_results.pivot(index='Max Depth', columns='Learning Rate',
    ↪values='Test Accuracy')
sns.heatmap(pivot_test, annot=True, fmt=".4f", cmap='viridis', ax=ax[0])
ax[0].set_title('TEST Accuracy (Higher is Better)')

# Heatmap 2: Overfitting Gap (What we want to minimize)
pivot_gap = df_results.pivot(index='Max Depth', columns='Learning Rate',
    ↪values='Gap')
sns.heatmap(pivot_gap, annot=True, fmt=".4f", cmap='coolwarm', ax=ax[1])
ax[1].set_title('Overfitting GAP (Lower is Better)')

plt.show()


# 4. Auto-select the winner (Based on best TEST accuracy)
best_row = df_results.loc[df_results['Test Accuracy'].idxmax()]
print("\n WINNER PARAMETERS:")
print(f"Best Depth: {int(best_row['Max Depth'])}")
print(f"Best Learning Rate: {best_row['Learning Rate']}")
print(f"Test Accuracy: {best_row['Test Accuracy']*100:.2f}%")
print(f"Train Accuracy: {best_row['Train Accuracy']*100:.2f}% (Gap:
    ↪{best_row['Gap']*100:.2f}%)")

```

Starting Deep Dive Grid Search (Train vs Test)...

Depth	Rate	Train Acc	Test Acc	Gap (Overfit)
10	0.01	95.62%	95.48%	0.15%
10	0.05	96.34%	96.07%	0.27%
10	0.1	96.63%	96.29%	0.34%
10	0.15	96.78%	96.36%	0.42%
10	0.2	96.91%	96.41%	0.50%
15	0.01	96.56%	96.06%	0.50%
15	0.05	97.19%	96.38%	0.81%
15	0.1	97.36%	96.44%	0.92%
15	0.15	97.52%	96.46%	1.05%
15	0.2	97.71%	96.45%	1.26%
20	0.01	96.95%	96.18%	0.77%
20	0.05	97.76%	96.39%	1.37%
20	0.1	97.95%	96.42%	1.53%
20	0.15	98.09%	96.41%	1.68%
20	0.2	98.29%	96.39%	1.90%
25	0.01	97.06%	96.17%	0.89%
25	0.05	98.07%	96.34%	1.73%
25	0.1	98.30%	96.36%	1.94%
25	0.15	98.43%	96.34%	2.09%
25	0.2	98.58%	96.32%	2.27%
30	0.01	97.07%	96.16%	0.91%

30	0.05	98.20%	96.30%	1.90%
30	0.1	98.46%	96.32%	2.15%
30	0.15	98.61%	96.27%	2.33%
30	0.2	98.75%	96.24%	2.51%
35	0.01	97.07%	96.16%	0.91%
35	0.05	98.26%	96.27%	1.99%
35	0.1	98.57%	96.28%	2.28%
35	0.15	98.72%	96.25%	2.47%
35	0.2	98.84%	96.21%	2.63%



WINNER PARAMETERS:

Best Depth: 15

Best Learning Rate: 0.15

Test Accuracy: 96.46%

Train Accuracy: 97.52% (Gap: 1.05%)

```
[21]: import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt
import xgboost as xgb
from sklearn.model_selection import cross_validate # <--- The new tool

# 1. Search Ranges
param_grid = {
    'max_depth': [15],
    'learning_rate': [0.01, 0.05, 0.1, 0.15, 0.17, 0.2, 0.25, 0.3, 0.4]
}

results = []
```

```

print("Starting Deep Dive Grid Search (Train vs Test)...")
print("-----")
print(f"{'Depth':<6} | {'Rate':<6} | {'Train Acc':<10} | {'Test Acc':<10} |  

    ↳{'Gap (Overfit)':<12}")
print("-----")

# 2. Double Loop
for depth in param_grid['max_depth']:
    for rate in param_grid['learning_rate']:

        model = xgb.XGBClassifier(
            n_estimators=100,
            max_depth=depth,
            learning_rate=rate,
            eval_metric='logloss'
        )

        # Use cross_validate to get BOTH scores
        cv_results = cross_validate(
            model,
            X_train,
            y_train,
            cv=3,
            scoring='accuracy',
            return_train_score=True # <--- This is the magic key
        )

        # Calculate averages
        mean_train = cv_results['train_score'].mean()
        mean_test = cv_results['test_score'].mean()
        gap = mean_train - mean_test

        results.append([depth, rate, mean_train, mean_test, gap])

        # Print row-by-row
        print(f"{'depth':<6} | {'rate':<6} | {'mean_train*100:.2f}% |  

            ↳{'mean_test*100:.2f}% | {'gap*100:.2f}%")

print("-----")

# 3. Visualize Results
df_results = pd.DataFrame(results, columns=['Max Depth', 'Learning Rate',  

    ↳'Train Accuracy', 'Test Accuracy', 'Gap'])

# Create two heatmaps side-by-side
fig, ax = plt.subplots(1, 2, figsize=(16, 6))

```


```
# Heatmap 1: Test Accuracy (What we want to maximize)
pivot_test = df_results.pivot(index='Max Depth', columns='Learning Rate',
    ↪values='Test Accuracy')
sns.heatmap(pivot_test, annot=True, fmt=".4f", cmap='viridis', ax=ax[0])
ax[0].set_title('TEST Accuracy (Higher is Better)')

# Heatmap 2: Overfitting Gap (What we want to minimize)
pivot_gap = df_results.pivot(index='Max Depth', columns='Learning Rate',
    ↪values='Gap')
sns.heatmap(pivot_gap, annot=True, fmt=".4f", cmap='coolwarm', ax=ax[1])
ax[1].set_title('Overfitting GAP (Lower is Better)')

plt.show()
```

Starting Deep Dive Grid Search (Train vs Test)...

Depth	Rate	Train Acc	Test Acc	Gap (Overfit)
15	0.01	96.56%	96.06%	0.50%
15	0.05	97.19%	96.38%	0.81%
15	0.1	97.36%	96.44%	0.92%
15	0.15	97.52%	96.46%	1.05%
15	0.17	97.59%	96.45%	1.14%
15	0.2	97.71%	96.45%	1.26%
15	0.25	97.86%	96.44%	1.42%
15	0.3	98.01%	96.42%	1.58%
15	0.4	98.29%	96.36%	1.93%



```
[22]: # 1. Retrieve the Winning Parameters from the previous step
# We access the 'best_row' variable we created in the Grid Search cell
final_depth = 15
final_rate = 0.15

print(f" Initializing Final Model with Winner Params: Depth={final_depth},
      ↪Rate={final_rate}")

# 2. Train the Final Model on ALL training data
final_model = xgb.XGBClassifier(
    n_estimators=200,          # We assume 200 is good for final polish
    max_depth=final_depth,    # The Winner
    learning_rate=final_rate, # The Winner
    eval_metric='logloss'
)

final_model.fit(X_train, y_train)

# 3. Final Evaluation on Test Set
# This is the "Moment of Truth" - checking on data the model has NEVER seen
y_pred_final = final_model.predict(X_test)
final_acc = accuracy_score(y_test, y_pred_final)

print(f"\n=====")
print(f" FINAL MODEL ACCURACY: {final_acc * 100:.2f}%")
print(f"=====")


# 4. Feature Importance
# See which new features (like 'dots' or 'keywords') actually mattered
plt.figure(figsize=(10, 8))
xgb.plot_importance(final_model, max_num_features=15, height=0.5,
    ↪importance_type='weight', title='Feature Importance (What signals phishing?
    ↪)')
plt.show()

# 5. Save the Brain
final_model.save_model("enhanced_url_classifier.json")
print("Model saved successfully.")
```

Initializing Final Model with Winner Params: Depth=15, Rate=0.15

```
=====
FINAL MODEL ACCURACY: 96.59%
=====
```

<Figure size 1000x800 with 0 Axes>



Model saved successfully.

[]: